

# Atividade 1 - Pré-processamento (Transformação dos dados) - base de dados **Hotel Reservations Kaggle**

## Análise Descritiva

### 1. Visualização de Dados

- Apresentar um mapa de calor (heatmap) entre todas as variáveis numéricas.

### 2. Transformação de Dados

- Efetuar as devidas transformações nos atributos categóricos.

### 3. Normalização de Dados

- Normalizar por Z-Score (Standard Scaler).

## Análise Agrupada

### 4. Group-by

- Utilizando Group-by, responder as seguintes perguntas:
  - A. Apresentar os valores mínimo, máximo e média do preço das diárias ( avg\_price\_per\_room ) agrupados por tipo de quarto ( room\_type\_reserved ).
  - B. Apresentar o valor médio de adultos e crianças hospedados em 2017 e 2018.

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# We want our plots to appear in the notebook
%matplotlib inline

from sklearn.preprocessing import LabelEncoder

## Models
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier

from sklearn.feature_selection import SelectKBest, mutual_info_classif, f
from sklearn.feature_selection import RFE, SequentialFeatureSelector as S
from feature_engine.selection import DropCorrelatedFeatures
from sklearn.pipeline import Pipeline
## Model evaluators
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.model_selection import RandomizedSearchCV, GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report
```

```
from sklearn.metrics import precision_score, accuracy_score, recall_score
```

```
In [ ]: df = pd.read_csv('Hotel Reservations.csv')
df.head(5)
```

```
Out[ ]:
```

	Booking_ID	no_of_adults	no_of_children	no_of_weekend_nights	no_of_week_night
0	INN00001	2	0	1	
1	INN00002	2	0	2	
2	INN00003	1	0	2	
3	INN00004	2	0	0	
4	INN00005	2	0	1	

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 36275 entries, 0 to 36274
```

```
Data columns (total 19 columns):
```

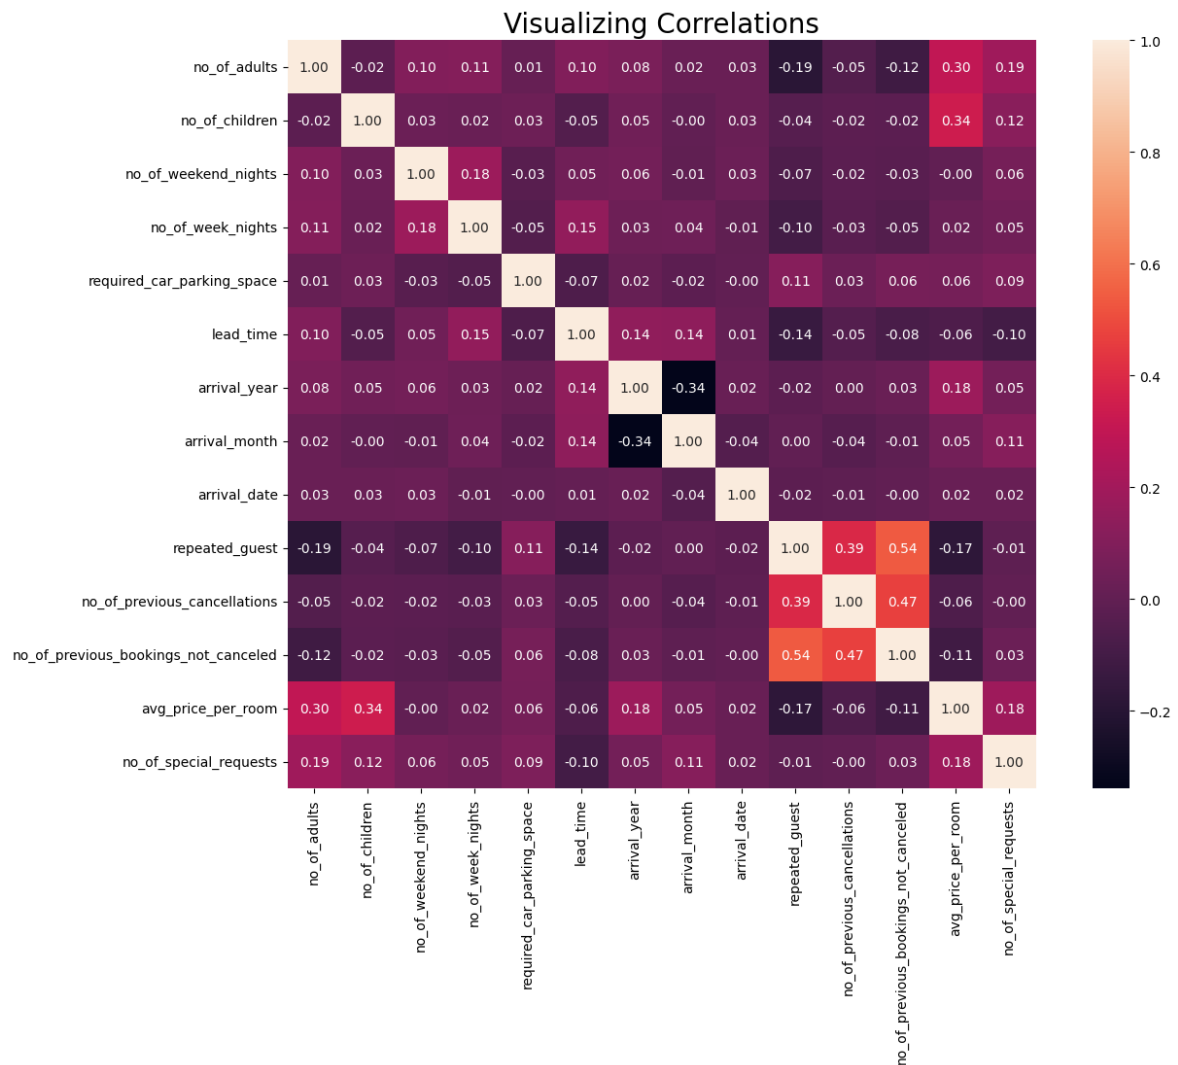
#	Column	Non-Null Count	Dtype
0	Booking_ID	36275 non-null	object
1	no_of_adults	36275 non-null	int64
2	no_of_children	36275 non-null	int64
3	no_of_weekend_nights	36275 non-null	int64
4	no_of_week_nights	36275 non-null	int64
5	type_of_meal_plan	36275 non-null	object
6	required_car_parking_space	36275 non-null	int64
7	room_type_reserved	36275 non-null	object
8	lead_time	36275 non-null	int64
9	arrival_year	36275 non-null	int64
10	arrival_month	36275 non-null	int64
11	arrival_date	36275 non-null	int64
12	market_segment_type	36275 non-null	object
13	repeated_guest	36275 non-null	int64
14	no_of_previous_cancellations	36275 non-null	int64
15	no_of_previous_bookings_not_canceled	36275 non-null	int64
16	avg_price_per_room	36275 non-null	float64
17	no_of_special_requests	36275 non-null	int64
18	booking_status	36275 non-null	object

```
dtypes: float64(1), int64(13), object(5)
```

```
memory usage: 5.3+ MB
```

## 1. Apresentar um mapa de calor (heatmap) entre todas as variáveis numéricas

```
In [ ]: plt.figure(figsize=(15, 10))
sns.heatmap(df.select_dtypes(include=['int64', 'float64']).corr(), square
plt.title("Visualizing Correlations", size = 20)
plt.show()
```



**4.A Apresentar os valores mínimo, máximo e média do preço das diárias avg\_price\_per\_room agrupados por tipo de quarto room\_type\_reserved.**

```
In [ ]: # Calculate the minimum, maximum, and average price of the daily rates gr
grouped_data = df.groupby('room_type_reserved')['avg_price_per_room'].agg


# Rename columns for clarity
grouped_data.columns = ['Room Type', 'Minimum Price', 'Maximum Price', 'A
grouped_data
```

Out[ ]:

	Room Type	Minimum Price	Maximum Price	Average Price
0	Room_Type 1	0.0	540.00	95.918532
1	Room_Type 2	0.0	284.10	87.848555
2	Room_Type 3	0.0	130.00	73.678571
3	Room_Type 4	0.0	375.50	125.287317
4	Room_Type 5	0.0	250.00	123.733623
5	Room_Type 6	0.0	349.63	182.212836
6	Room_Type 7	0.0	306.00	155.198291

```
In [ ]: room_type_prices = df.groupby('room_type_reserved')['avg_price_per_room']

room_type_prices.plot(kind='bar', figsize=(8, 6),)
plt.xlabel('Room Type')
plt.ylabel('Average Daily Rate (ADR)')
plt.show()
```



## 4.2 Apresentar os valor médio de adultos e crianças hospedados em 2017 e 2018.

```
In [ ]: # Agrupar por 'arrival_year' e calcular a média de 'no_of_adults' e 'no_o
```

```
grouped_data = df.groupby('arrival_year')[['no_of_adults', 'no_of_childre  
  
# Renomear colunas para clareza  
grouped_data.columns = ['Ano de Chegada', 'Média de Adultos', 'Média de C  
  
grouped_data
```

```
Out[ ]:      Ano de Chegada  Média de Adultos  Média de Crianças  
0          2017          1.759902          0.065705  
1          2018          1.863580          0.113941
```

## 2. Efetuar as devidas transformações nos atributos categóricos

```
In [ ]: df.dtypes
```

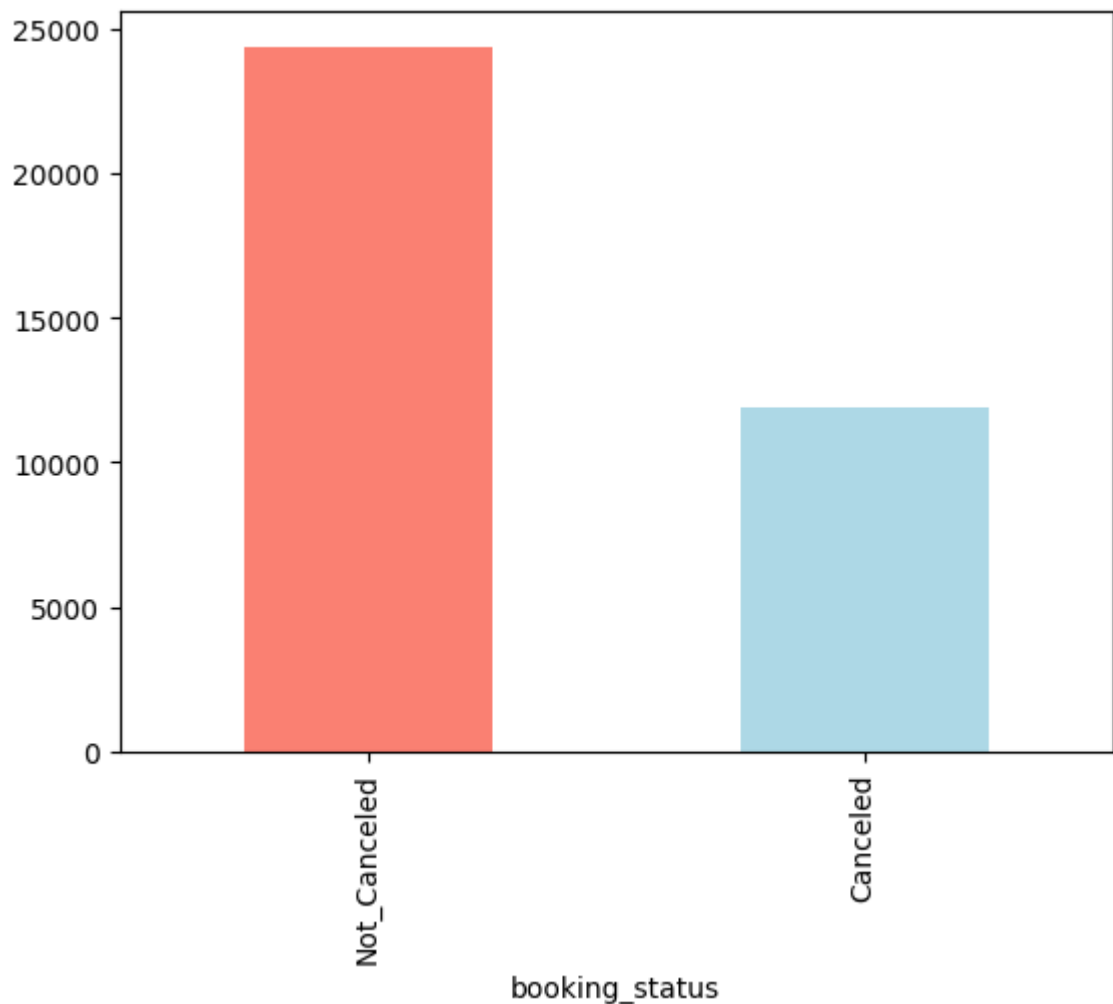
```
Out[ ]: Booking_ID          object  
no_of_adults          int64  
no_of_children          int64  
no_of_weekend_nights    int64  
no_of_week_nights      int64  
type_of_meal_plan      object  
required_car_parking_space  int64  
room_type_reserved      object  
lead_time              int64  
arrival_year           int64  
arrival_month           int64  
arrival_date           int64  
market_segment_type     object  
repeated_guest          int64  
no_of_previous_cancellations  int64  
no_of_previous_bookings_not_canceled  int64  
avg_price_per_room      float64  
no_of_special_requests  int64  
booking_status          object  
dtype: object
```

## pré-processamento

```
In [ ]: df = df.drop('Booking_ID', axis = 1)  
df = df.drop('arrival_year', axis = 1)
```

Vou dropar o Booking ID e o ano da reserva para evitar overfitting, já que o ID poderia acabar indicando uma relação entre as reservas e o ano é uma feature que não vai se repetir nos anos seguintes, não sendo útil para a identificação de cancelamentos futuros

```
In [ ]: # Target variable  
df.booking_status.value_counts().plot(kind="bar", color=["salmon", "light
```



```
In [ ]: # Import label encoder
        from sklearn import preprocessing

        # label_encoder object knows how to understand word labels.
        label_encoder = preprocessing.LabelEncoder()

        # Encode labels in column 'species'.
        df['booking_status'] = label_encoder.fit_transform(df['booking_status'])

        df['booking_status'].unique()
```

```
Out[ ]: array([1, 0])
```

```
In [ ]: # Make a copy of the original DataFrame to perform edits on
        #df_tmp = df.copy()
```

```
In [ ]: #names of columns
        columns = list(df.columns)
        columns
```

```
Out[ ]: ['no_of_adults',
        'no_of_children',
        'no_of_weekend_nights',
        'no_of_week_nights',
        'type_of_meal_plan',
        'required_car_parking_space',
        'room_type_reserved',
        'lead_time',
        'arrival_month',
        'arrival_date',
        'market_segment_type',
        'repeated_guest',
        'no_of_previous_cancellations',
        'no_of_previous_bookings_not_canceled',
        'avg_price_per_room',
        'no_of_special_requests',
        'booking_status']
```

```
In [ ]: # Check for missing values
        #df_tmp.isna().sum()
```

## Analisar as variáveis to encode

```
In [ ]: categorical_columns = df.select_dtypes(include=['object']).columns
```

```
In [ ]: for categorical_feature in categorical_columns:
        print(f'{categorical_feature}: {df[categorical_feature].unique()}')
```

```
type_of_meal_plan: ['Meal Plan 1' 'Not Selected' 'Meal Plan 2' 'Meal Plan
3']
room_type_reserved: ['Room_Type 1' 'Room_Type 4' 'Room_Type 2' 'Room_Type
6' 'Room_Type 5'
'Room_Type 7' 'Room_Type 3']
market_segment_type: ['Offline' 'Online' 'Corporate' 'Aviation' 'Complemen
tary']
```

```
In [ ]: columns_to_encode = [ 'market_segment_type', 'type_of_meal_plan', 'room_t
encoder = LabelEncoder()
for column in columns_to_encode:
    df[column] = encoder.fit_transform(df[column])

# Applying one-hot encoding
#df = pd.get_dummies(df, columns=['type_of_meal_plan', 'room_type_reserve
```

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 36275 entries, 0 to 36274
Data columns (total 17 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0   no_of_adults                                  36275 non-null   int64
1   no_of_children                                36275 non-null   int64
2   no_of_weekend_nights                          36275 non-null   int64
3   no_of_week_nights                             36275 non-null   int64
4   type_of_meal_plan                             36275 non-null   int64
5   required_car_parking_space                    36275 non-null   int64
6   room_type_reserved                            36275 non-null   int64
7   lead_time                                     36275 non-null   int64
8   arrival_month                                 36275 non-null   int64
9   arrival_date                                  36275 non-null   int64
10  market_segment_type                           36275 non-null   int64
11  repeated_guest                                36275 non-null   int64
12  no_of_previous_cancellations                  36275 non-null   int64
13  no_of_previous_bookings_not_canceled          36275 non-null   int64
14  avg_price_per_room                            36275 non-null   float64
15  no_of_special_requests                        36275 non-null   int64
16  booking_status                                36275 non-null   int64
dtypes: float64(1), int64(16)
memory usage: 4.7 MB
```

```
In [ ]: df.head(10)
```

```
Out[ ]:   no_of_adults  no_of_children  no_of_weekend_nights  no_of_week_nights  type_of_r
```

0	2	0	1	2
1	2	0	2	3
2	1	0	2	1
3	2	0	0	2
4	2	0	1	1
5	2	0	0	2
6	2	0	1	3
7	2	0	1	3
8	3	0	0	4
9	2	0	0	5

## Rescaling the data

```
In [ ]: from sklearn.preprocessing import StandardScaler

scaler_cols=['no_of_adults',
             'no_of_children',
             'no_of_weekend_nights',
             'no_of_week_nights',
             'required_car_parking_space',
             'lead_time',
             'arrival_month',
```



```

    'arrival_date',
    'repeated_guest',
    'no_of_previous_cancellations',
    'no_of_previous_bookings_not_canceled',
    'avg_price_per_room',
    'no_of_special_requests']
scaler = StandardScaler()
ajuste = scaler.fit(df[scaler_cols])
df[scaler_cols] = ajuste.transform(df[scaler_cols])

```

```
In [ ]: df.head()
```

```
Out[ ]:
```

	no_of_adults	no_of_children	no_of_weekend_nights	no_of_week_nights	type_of_r
0	0.298893	-0.26147	0.217401	-0.144803	
1	0.298893	-0.26147	1.365993	0.563972	
2	-1.628975	-0.26147	1.365993	-0.853578	
3	0.298893	-0.26147	-0.931190	-0.144803	
4	0.298893	-0.26147	0.217401	-0.853578	

```
In [ ]: #plt.figure(figsize=(20, 20))
#sns.heatmap(df.select_dtypes(include=['int64', 'float64']).corr(), squar
#plt.title("Visualizing Correlations", size = 20)
#plt.show()
```

## Optional

```
In [ ]: # Everything except target variable
X = df.drop("booking_status", axis=1)

# Target variable
y = df['booking_status']
```

## train and test split

```
In [ ]: # Random seed for reproducibility
np.random.seed(55)

# Split into train & test set
X_train, X_test, y_train, y_test = train_test_split(X, # independent vari
y, # dependent variab
test_size = 0.2) # pe
```

## Model choices

We'll be using the following and comparing their results.

1. Decision Tree - DecisionTreeClassifier()
2. K-Nearest Neighbors - KNeighborsClassifier()
3. RandomForest - RandomForestClassifier()

```
In [ ]: models = {  
    "Decision Tree": DecisionTreeClassifier(),  
    "KNN": KNeighborsClassifier(),  
    "Random Forest": RandomForestClassifier()  
}
```

```
In [ ]: from sklearn.metrics import accuracy_score, recall_score, precision_score  
  
def fit_and_evaluate(models, X_train, X_test, y_train, y_test):  
    """  
    Fits and evaluates given machine learning models.  
    models: a dict of different Scikit-Learn machine learning models  
    X_train: training data  
    X_test: testing data  
    y_train: target training data  
    y_test: target test data  
    """  
  
    np.random.seed(42) # Random seed for reproducible results  
    model_metrics = {} # Dictionary to keep all model metrics  
  
    # Loop through models  
    for name, model in models.items():  
        model.fit(X_train, y_train) # Fit the model to the data  
        y_pred = model.predict(X_test) # Predictions  
        # Calculate metrics  
        accuracy = accuracy_score(y_test, y_pred)  
        recall = recall_score(y_test, y_pred)  
        precision = precision_score(y_test, y_pred)  
        f1 = f1_score(y_test, y_pred)  
  
        # Store metrics in the dictionary  
        model_metrics[name] = {  
            "Accuracy": accuracy,  
            "Recall": recall,  
            "Precision": precision,  
            "F1 Score": f1  
        }  
  
    return model_metrics
```

## Model comparison

```
In [ ]: # Call the function  
model_metrics = fit_and_evaluate(models=models,  
                                X_train=X_train,  
                                X_test=X_test,  
                                y_train=y_train,
```

```

y_test=y_test)

# Convert model metrics to DataFrame for better visualization
model_metrics_df = pd.DataFrame(model_metrics).T
# Display the DataFrame
print(model_metrics_df)

```

	Accuracy	Recall	Precision	F1 Score
Decision Tree	0.867815	0.890960	0.914609	0.902630
KNN	0.854859	0.902385	0.888319	0.895297
Random Forest	0.905720	0.943275	0.921480	0.932250

## feature selection

```

In [ ]: # Definindo os classificadores
classifiers = {
    "Decision Tree": DecisionTreeClassifier(),
    "KNN": KNeighborsClassifier(),
    "Random Forest": RandomForestClassifier()
}

# Técnicas de seleção de características
feature_selection_techniques = {
    "Mutual Information": SelectKBest(mutual_info_classif, k=9),
    "ANOVA": SelectKBest(f_classif, k=9),
    "Drop Correlated": DropCorrelatedFeatures(threshold=0.8, method='pear')
    # Add SmartCorrelatedGroups após
}

## Configuração do pipeline de avaliação
def evaluate_pipeline(X_train, X_test, y_train, y_test, classifier, selector):
    pipeline = Pipeline([
        ('selector', selector),
        ('classifier', classifier)
    ])

    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)

    return {
        "Accuracy": accuracy_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred, average='binary'),
        "Precision": precision_score(y_test, y_pred, average='binary'),
        "F1 Score": f1_score(y_test, y_pred, average='binary')
    }

```

```

In [ ]: results = {}

for clf_name, clf in classifiers.items():
    for fs_name, fs in feature_selection_techniques.items():
        key = f"{clf_name} + {fs_name}"
        results[key] = evaluate_pipeline(X_train, X_test, y_train, y_test)

# Convertendo resultados em DataFrame para visualização
results_df = pd.DataFrame(results).T
print(results_df)

```

	Accuracy	Recall	Precision	F1 Score
e				
Decision Tree + Mutual Information	0.863818	0.887753	0.911880	0.89965
5				
Decision Tree + ANOVA	0.853480	0.880136	0.904242	0.89202
6				
Decision Tree + Drop Correlated	0.870986	0.893766	0.916547	0.90501
3				
KNN + Mutual Information	0.862991	0.905191	0.896565	0.90085
8				
KNN + ANOVA	0.856099	0.908398	0.885329	0.89671
5				
KNN + Drop Correlated	0.854859	0.902385	0.888319	0.89529
7				
Random Forest + Mutual Information	0.889318	0.922229	0.917265	0.91974
0				
Random Forest + ANOVA	0.885596	0.924434	0.910563	0.91744
6				
Random Forest + Drop Correlated	0.905720	0.942874	0.921811	0.93222
4				

## Optional2 only for studies purpose

Hyperparameter tuning and cross-validation Next steps to be taken:

1. Tune model hyperparameters
2. Perform cross-validation
3. Plot ROC curves
4. Make a confusion matrix
5. Get precision, recall and F1-score metrics
6. Find the most important model features

## 1. Tune model hyperparameters

Tune by hand

```
In [ ]: train_scores = []

test_scores = []

# Create a list of different values for n_neighbors
neighbors = range(1, 11) # 1 to 10

# Setup algorithm
knn = KNeighborsClassifier()

# Loop through different neighbors values
for i in neighbors:
    knn.set_params(n_neighbors = i) # set neighbors value

    # Fit the algorithm
    knn.fit(X_train, y_train)


    # Update the training scores
    train_scores.append(knn.score(X_train, y_train))
```

```
# Update the test scores
test_scores.append(knn.score(X_test, y_test))
```

```
In [ ]: #visualizing the scores
plt.plot(neighbors, train_scores, label="Train score")
plt.plot(neighbors, test_scores, label="Test score")
plt.xticks(np.arange(1, 11, 1))
plt.xlabel("Number of neighbors")
plt.ylabel("Model score")
plt.legend()

print(f"Maximum KNN score on the test data: {max(test_scores)*100:.2f}%")
```

Maximum KNN score on the test data: 85.49%



- Tuning models with with RandomizedSearchCV

```
In [ ]: # Different LogisticRegression hyperparameters
#log_reg_grid = {"C": np.logspace(-4, 4, 20),
#               "solver": ["liblinear"]}

rf_grid = {
    'n_estimators': np.arange(10, 100, 3), # Number of trees in the fore
    'max_depth': [None] + list(np.arange(5, 30)), # Maximum number of le
    'min_samples_split': np.arange(2, 20), # Minimum number of samples r
    'min_samples_leaf': np.arange(1, 20), # Minimum number of samples re
    'max_features': [ 'sqrt', 'log2'], # Number of features to consider
    'bootstrap': [True, False] # Method of selecting samples for trainin
}
```

Now we tune RandomForestClassifier using RandomizedSearchCV

```
In [ ]: # Setup random seed
np.random.seed(90)

# Setup random hyperparameter search for RandomForestClassifier
rs_rf = RandomizedSearchCV(RandomForestClassifier(),
                           param_distributions=rf_grid,
                           cv=5,
                           n_iter=20,
                           verbose=True)

# Fit random hyperparameter search model
rs_rf.fit(X_train, y_train)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

```
Out[ ]: ► RandomizedSearchCV ⓘ ?
        ► estimator: RandomForestClassifier
          ► RandomForestClassifier ?
```

```
In [ ]: # Find the best parameters
rs_rf.best_params_
```

```
Out[ ]: {'n_estimators': 79,
        'min_samples_split': 2,
        'min_samples_leaf': 1,
        'max_features': 'log2',
        'max_depth': 24,
        'bootstrap': False}
```

```
In [ ]: # Evaluate the randomized search random forest model acc
rs_rf.score(X_test, y_test)
```

```
Out[ ]: 0.9046175051688491
```

- Tuning a model with GridSearchCV
- Evaluating a classification model

We want:

1. ROC curve and AUC score - `plot_roc_curve()`
2. Confusion matrix - `confusion_matrix()`
3. Classification report - `classification_report()`
4. Precision - `precision_score()`
5. Recall - `recall_score()`
6. F1-score - `f1_score()`

```
In [ ]: # Make preidctions on test data
y_preds = rs_rf.predict(X_test)
```


```
In [ ]: y_preds
```

```
Out[ ]: array([1, 1, 0, ..., 1, 1, 1])
```

```
In [ ]: # Import ROC curve function from metrics module
        from sklearn.metrics import RocCurveDisplay

        # Plot ROC curve and calculate AUC metric
        RocCurveDisplay.from_estimator(rs_rf, X_test, y_test)
```

```
Out[ ]: <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7f28bcb04b50>
```



```
In [ ]: from sklearn.metrics import roc_auc_score

        roc_auc_score(y_test, y_preds)
```

```
Out[ ]: 0.8816313365131445
```

- confusion matrix

```
In [ ]: # Display confusion matrix
        print(confusion_matrix(y_test, y_preds))
```

```
[[1859  407]
 [ 285 4704]]
```

```
In [ ]: # Import Seaborn
        import seaborn as sns
        sns.set(font_scale=1.5) # Increase font size


        def plot_conf_mat(y_test, y_preds):
            """
            Plots a confusion matrix using Seaborn's heatmap().
```

```

"""
fig, ax = plt.subplots(figsize=(3, 3))
ax = sns.heatmap(confusion_matrix(y_test, y_preds),
                  annot=True, # Annotate the boxes
                  cbar=False)
plt.xlabel("true label")
plt.ylabel("predicted label")

plot_conf_mat(y_test, y_preds)

```



```

In [ ]: # Show classification report
print(classification_report(y_test, y_preds))

```

	precision	recall	f1-score	support
0	0.87	0.82	0.84	2266
1	0.92	0.94	0.93	4989
accuracy			0.90	7255
macro avg	0.89	0.88	0.89	7255
weighted avg	0.90	0.90	0.90	7255

```

In [ ]: # Check best hyperparameters
rs_rf.best_params_

```

```

Out[ ]: {'n_estimators': 79,
        'min_samples_split': 2,
        'min_samples_leaf': 1,
        'max_features': 'log2',
        'max_depth': 24,
        'bootstrap': False}

```

```

In [ ]:

```